

# Lung Nodule Segmentation — DynUNet / 3D U-Net (v9)

---

## 1. Task

---

Voxel-level segmentation of pulmonary nodules in 3D chest CT volumes, cropped to a per-series lung bounding box.

- **Input** (1, 256, 256, 256) CT crop, intensity-normalised to [0, 1], resampled from the bbox (padding = 20 vox).
- **Output** (2, 256, 256, 256) softmax logits — class 0 = background, class 1 = nodule.

## 2. Model

---

[MONAI DynUNet](#) — nnU-Net-style **3D U-Net**, DynUNet (nnU-Net-style 3D U-Net).

| Property             | Value                            |
|----------------------|----------------------------------|
| Architecture family  | DynUNet (nnU-Net-style 3D U-Net) |
| Spatial dims         | 3                                |
| Input channels       | 1                                |
| Output channels      | 2 (softmax)                      |
| Encoder levels       | 6                                |
| Kernel sizes         | [3, 3, 3] at every level         |
| Strides              | [[1,1,1], [2,2,2]×5]             |
| Normalisation        | instance                         |
| Deep supervision     | off                              |
| Trainable parameters | <b>31,181,730</b>                |

## 3. Data

---

Trained on the `unified_v2` split (patient-grouped, dataset-stratified, **full corpus**): NLST + NSCLC + LIDC-IDRI.

Split sizes: **1 683 train / 297 val / 325 test (held out)**.

Class imbalance: nodule voxels are on the order of  $\sim 10^{-5}$  of the total, which shapes how each metric should be read (§7).

## 4. Training

| Setting         | Value                                                                                                                               |
|-----------------|-------------------------------------------------------------------------------------------------------------------------------------|
| Loss            | Focal Tversky + weighted CE ( $\alpha=0.3$ , $\beta=0.7$ , $\gamma=2.0$ , $\lambda_{CE}=0.1$ , <code>ce_nodule_weight</code> = 100) |
| Optimizer       | Adam                                                                                                                                |
| Learning rate   | $1 \times 10^{-5}$                                                                                                                  |
| Weight decay    | $1 \times 10^{-5}$                                                                                                                  |
| LR schedule     | Cosine annealing, $T_{max} = 400$ , $\eta_{min} = 1 \times 10^{-6}$                                                                 |
| Batch size      | 2                                                                                                                                   |
| Epochs          | 400                                                                                                                                 |
| Random seed     | 42                                                                                                                                  |
| Mixed precision | bf16 (autocast, no GradScaler)                                                                                                      |
| Augmentation    | 3D flips, 90° rotations, elastic rotation, zoom, intensity scale/shift, Gaussian noise/blur, contrast                               |
| Hardware        | 1 × NVIDIA H100 94 GB                                                                                                               |
| Wall-clock      | ≈ 3.5 days                                                                                                                          |

Best checkpoint (evaluated below): **epoch 354 / 400**.

## 5. Evaluation protocol

**Held-out test split** (325 series). The training loop used the `val` split of `unified_v2` for early stopping and best-Dice checkpoint selection; the numbers below come from the completely untouched **test** split (never seen during training or model selection).

Predictions taken as `argmax` over the 2-channel softmax output; foreground = class 1. All metrics are computed at the voxel level and micro-averaged over the whole test split (5,452,595,200 voxels total).

## 6. Results

### Ticket-required metrics

| Metric           | Value         |
|------------------|---------------|
| <b>mean IoU</b>  | <b>0.7594</b> |
| <b>Accuracy</b>  | <b>0.9993</b> |
| <b>Precision</b> | <b>0.6236</b> |
| <b>Recall</b>    | <b>0.7569</b> |

## Supplementary

| Metric                          | Value  |
|---------------------------------|--------|
| IoU (foreground / nodule class) | 0.5196 |
| IoU (background class)          | 0.9993 |
| Dice / F1 (micro)               | 0.6838 |

### Per-case Dice distribution (325 val cases)

| Statistic | Value  |
|-----------|--------|
| Mean      | 0.5560 |
| Std       | 0.2551 |
| Min       | 0.0000 |
| p05       | 0.0031 |
| p25       | 0.3837 |
| Median    | 0.6196 |
| p75       | 0.7621 |
| p95       | 0.8709 |
| Max       | 0.9455 |

## 7. Notes on interpretation

Nodule segmentation is severely class-imbalanced ( $\sim 10^{-5}$  of voxels are nodule), which distorts the standard metric set:

- **Accuracy** is trivially near 1.0 for any reasonable model — the model gets  $\approx 99.9313$  % of voxels right by predicting "background" almost everywhere. Uninformative on its own here.
- **mIoU** averages the foreground and background IoU. `IoU_background` is essentially 1.0, so mIoU is roughly  $0.5 + 0.5 \cdot \text{IoU\_nodule}$ . The useful signal is in `IoU_nodule` (0.5196) and Dice / F1 (0.6838).
- **Precision / Recall** are the standard per-voxel figures — no imbalance caveat needed. Recall 0.7569 means the model correctly labels  $\sim 76$  % of nodule voxels; Precision 0.6236 means  $\sim 62$  % of predicted-nodule voxels are true nodule.

Per-case Dice mean (0.5560) is lower than the micro Dice (0.6838) because large nodules dominate the micro average; per-case Dice weights each patient equally regardless of nodule size. Clinically the per-case distribution is the more useful summary.

## 8. How to load & run inference

---

Pre-trained weights are hosted at [Kakimaki00/nodule-dynunet-3d](#) . Download and load directly:

```
huggingface-cli download Kakimaki00/nodule-dynunet-3d --local-dir ./ckpt
```

```
import yaml, torch
from monai.networks.nets import DynUNet

cfg = yaml.safe_load(open("config.yaml"))["model"]
default_strides = [[1, 1, 1]] + [[2, 2, 2]] * 5
strides = cfg.get("strides", default_strides)
kernel = cfg.get("kernel_size", [[3, 3, 3]] * len(strides))
model = DynUNet(
    spatial_dims      = cfg["spatial_dims"],
    in_channels       = cfg["in_channels"],
    out_channels      = cfg["out_channels"],
    kernel_size       = kernel,
    strides           = strides,
    upsample_kernel_size = cfg.get("upsample_kernel_size", strides[1:]),
    norm_name         = cfg.get("norm_name", "instance"),
    deep_supervision  = cfg.get("deep_supervision", False),
)
state = torch.load("model.pth", map_location="cpu", weights_only=True)
model.load_state_dict(state)
model.eval()

with torch.no_grad():
    # Input: (B, 1, 256, 256, 256), lung-bbox-cropped CT resampled to 256³.
    x = torch.randn(1, 1, 256, 256, 256)
    logits = model(x)  # (B, 2, D, H, W)
    pred_class = logits.argmax(dim=1)  # (B, D, H, W) in {0, 1}
    nodule_mask = (pred_class == 1).to(torch.uint8)
```

This model expects **lung-bbox-cropped** input. A separate 2D ROI model is needed to produce that bbox — see the accompanying ROI checkpoints ( [Kakimaki00/roi-segresnet-2d](#) or [Kakimaki00/roi-swinunetr-2d](#) ), or use the demo's `inference.py` which chains them for you.

## 9. Reproducibility

---

- Config `configs/v9.yaml`
- Checkpoint `checkpoints/v9/best_model.pth`
- Metrics JSON `checkpoints/v9/eval_metrics_nodule_test.json`
- Command `python eval_metrics_nodule.py --config configs/v9.yaml --ckpt checkpoints/v9/best_model.pth --split test`

Companion two-stage nodule models on HuggingFace:

- [Kakimaki00/nodule-segresnet-3d-small](#) — small SegResNet (v6)
- [Kakimaki00/nodule-segresnet-3d-wide](#) — wider SegResNet (paper's best)